

## Ansible Playbook Process:

### 1. Create Playbook

- Create an empty playbook

### 2. Outline each Task

- Name each task required

### 3. Build each Task

- Build each task required according to documentation

### 4. Check Playbook

- `ansible-playbook -C <playbook> -K`

## Critical Exam Tips:

- Always number your playbooks for speed, ex. **one, two three** (can't be numbers)
- Always use **^string.\*** when using regexp flag to make sure entire line gets replaced
- Always use **aliases** in `~/.bashrc` for `ansible-playbook` and `ansible-doc` (ap, and ad)
- Always use **3 slashes in baseurl** for FTP and HTTP
  - `ftp:///controller/(BaseOS/AppStream)` OR `http:///controller/(BaseOS/AppStream)`
- Always use **block/rescue/always** for error handling
- Use `---` and `import_playbook` with nothing else when using multiple playbooks in one file

## Critical Ad-Hoc Flags:

Description	Flag
Ask Become Pass	-K
Ask Pass	-k

Become	-b
Module Name	-m
Arguments	-a
Check syntax	-c
Inventory file	-I
Remote user	-u
Extra vars	-e

## Basics

Remote user must exist on remote host

Remote user must be in wheel group on remote host

Remote user must be able to do password-less sudo on remote host

SSH Public Key must be copied into ~/.ssh/authorized\_keys file on remote host under remote user home directory

## Critical Commands:

Description	Command
Display all Ansible keywords	ansible-doc --type keyword --list
Get help using a Module	ansible-doc <module>
Display all installed Ansible modules	ansible-doc -l
Display parameters for a Module	ansible-doc -s <module>
Test Reachability of Hosts – Ping	ansible all -m ping

Run Ansible Ad-Hoc commands	<code>ansible -m &lt;module&gt; -a &lt;commandarguments&gt;</code>
Which config file is in use	<code>ansible --version</code>
Which hosts are in inventory	<code>ansible all --list-hosts</code>
Modules related to gathering facts	<code>ansible-doc -l   grep fact</code>
List all gathering facts variables	<code>ansible -m setup &lt;hostname&gt; &gt; facts.yml</code>
Install Containerized AAP	<code>ansible-playbook -i inventory-growth ansible.containerized_installer.install -K</code>
Verify playbook syntax	<code>ansible-lint &lt;playbook/inventory/ansible.cfg&gt;</code>

## Documentation

All of these can be found under Ansible Core documentation - This is all you need

Documentation

[Building Ansible inventories](#)

[Using Ansible command line tools](#)

[Using Ansible playbooks](#)

[Protecting sensitive data with Ansible vault](#)

[Using Ansible modules and plugins](#)

[Using Ansible collections](#)

[Ansible tips and tricks](#)

Reference & Appendices

[Frequently Asked Questions](#)

[Ansible Configuration Settings](#)

[Playbook Keywords](#)

Special Variables

Controlling how Ansible behaves: precedence rules

YAML Syntax

### Hyphen or No-Hyphen:

Vars is a **dictionary of key/value pairs**. Each key is just one property, not an “item in a sequence,” so **no hyphen**

- **Playbooks, roles, tasks, handlers -> lists** (start with -)
- **Vars and module args -> mappings** (just key: value)

Playbooks will still run with vars missing hyphens, this is just good practice

### Playbook & Config Structure:

Refer to documentation on structure of both files

#### Playbooks

name

hosts

become / become\_user

vars / vars\_files

tasks

module

#### Config

[defaults]

inventory = inventory

remote\_user = awx

ask\_pass = false

```
private_key_file = /path/to/private_key
```

```
[privilege_escalation]
```

```
become = true
```

```
become_user = root
```

```
become_method = sudo
```

```
become_ask_pass = false
```

### **In project directory**

```
mkdir collections group_vars host_vars vars roles
```

### **Ansible Cheat Sheet:**

### **Ansible when:**

#### **Check if host is in a group**

```
- name: Install nmap
```

```
  dnf:
```

```
    name: nmap
```

```
    state: latest
```

```
  when: inventory_hostname in groups['dev']
```

#### **Check if host is not in a group**

```
- name: Install nmap
```

```
  dnf:
```

```
    name: nmap
```

```
    state: latest
```

```
  when: inventory_hostname not in groups['dev']
```

### **Check multiple conditions**

- name: Install nmap

dnf:

name: nmap

state: latest

when: "'dev' in group\_names or 'prod' in group\_names"

### **Check if a fact doesn't exist**

- name:

fail:

msg: "value is not in variable"

when: ansible\_devices.sdb is not defined

### **Check if a string equals something**

- name: Do something only on RHEL

command: echo "This is RHEL"

when: ansible\_distribution == "RedHat"

### **Check if boolean is true**

- name: Only run when debug\_mode is true

debug:

msg: "Debug mode enabled"

when: debug\_mode | bool

### **Check for value in a list**

- name: Only run for specific datacenter

debug:

msg: "This DC is allowed"

when: datacenter in ['dc1', 'dc2']

### **Check that a number starts with 1**

- name: Ensure users with UID starting with 1 are present

user:

```
state: present
name: "{{ item.username }}"
shell: /bin/bash
when:
  - item.uid | string is regex('^1')
loop: "{{ users }}"
```

### Check and

```
- name: Only run for specific datacenter
debug:
  msg: "This is true"
when: something==true and something2==true
```

### Check or

```
- name: Only run for specific datacenter
debug:
  msg: "This is true"
when: something==true or something2==true
```

## Variables

Inventory/group\_vars/host\_vars

### a) Inventory variables

- Set directly in the inventory file:

```
[web]
```

```
web1 ansible_host=192.168.1.10 username=julie
```

### b) Group variables (group\_vars/) -> Only defines variables for a specific group

- Define variables for a whole group of hosts.

- File structure:

```
group_vars/prod.yml
```

group\_vars/dev.yml

- Example group\_vars/prod.yml:

username: julie

- Automatically applied to hosts in that group.

### **c) Host variables (host\_vars/) -> Only defines variables for a specific host**

- Variables specific to one host.

- File structure:

host\_vars/web1.yml

- Example:

username: julie

timezone: UTC

## **Playbook/Task level**

### **a) vars: in a playbook**

- Example:

- hosts: prod

vars:

username: julie

tasks:

- debug:

msg: "User is {{ username }}"

### **b) vars: at task level**

- Example:

- debug:

msg: "User is {{ username }}"

vars:

username: julie

## **External files (vars\_files:)**



- Include a YAML file for variables:
- hosts: prod
- vars\_files:
  - vars/prod.yml
- tasks:
  - debug:
    - msg: "User is {{ username }}"

### Extra vars (CLI)

- Passed at runtime, highest precedence:
- ansible-playbook playbook.yml -e "username=julie"

### Facts

- Gathered automatically or set via set\_fact:
- set\_fact:
  - username: julie
- Can be host-specific and dynamic.

### Precedence Summary (RHCE-level)

From lowest → highest:

1. Inventory/group\_vars/host\_vars
2. Playbook vars:
3. vars\_files:
4. set\_fact
5. Extra vars (-e)

Extra vars always override everything else.

### Critical Information:

**Always use Debug module!!!**

**List all gathered facts**

ansible <hostname> -m setup > facts

## Use group vars for group-specific variables

### Create a custom fact

1. Create a file in /etc/ansible/facts.d on the managed host
2. Name the file ending in .fact
3. Put the following content in file

```
[fact_name]
```

```
custom_key=some_value
```

4. Verify the custom fact

```
ansible all -m setup -a "filter=ansible_local"
```

## This will set the facts every time ansible runs

If you want it to run once just use the **set\_fact** flags

- name: Example playbook

hosts: all

tasks:

- name: Set a custom fact

set\_fact:

my\_custom\_fact: "hello world"

- name: Show custom fact

debug:

msg: "My fact is {{ my\_custom\_fact }}"

## Using vars\_files

vars\_files:

- vars/user\_list.yml

You must create a vars directory

### **Use stat module to check if something exists**

stat -> register in a variable, use ansible-doc stat

variable.stat.exists and not variable.stat.exists

### **Use uri module for testing web service reachability**

uri and curl -I both return headers

### **Calling Facts (These are Equivalent)**

- debug:

msg: "{{ ansible\_facts['distribution\_release'] }}" -> **Remove ansible\_ from key name**

- debug:

msg: "{{ ansible\_distribution\_release }}" -> **Use whole key name**

### **Calling a Nested Fact**

- debug:

msg: "{{ ansible\_date\_time.date }}" -> **Calls a key within a Dictionary**

### **Create user on Remote Hosts**

ansible all -u awx -k -b -K -m user -a "name=remote\_user"

### **Set a password on user**

ansible all -u awx -k -b -K -m shell -a "echo <password> | passwd --stdin awx"

### **Create SSH keys, and Copy SSH keys**

for i in vm1, vm2; do ssh-copy-id \$i; done or sh-copy-id -i ~/.ssh/id\_rsa.pub awx@remote\_host

**For keywords consult:**

ansible-doc --type keyword -list

**Generate a Template ansible.cfg file (found in /etc/ansible/ansible.cfg)**

ansible-config init --disabled -t all > ansible.cfg

**Generate an Inventory file in yaml format (creates a yaml formatted inventory)**

ansible-inventory --list --yaml > inventory

**Enables privilege escalation, with sudo, no password prompt - Bad practice in production, good for lab**

ansible all -k -b -K -m copy -a "content='awx ALL=(ALL) NOPASSWD: ALL' dest=/etc/sudoers.d/awx"

**Update sudoers.d so Controller so user can escalate privileges without being prompted for a password**

echo "awx ALL=(ALL) NOPASSWD: ALL" >> /etc/sudoers.d/awx

**Copy everything stored on ISO to repo server starting with A or B (AppStream & BaseOS Repos)**

sudo cp -R /mnt/[AB]\* /reposerver

**Use command module to restore context after making changes using SELinux modules**

command: restorecon -Rv /localdirectory

**Fail if a command didn't successfully run:**

- name: ok

command: chronyc tracking || echo "Time could not be synchronized"

**Install ansible.posix (firewalld) module**

ansible-galaxy collection install ansible.posix

**Default directory cron module will create files in if cron\_file isn't used**

/var/spool/cron/

**Use Jinja2 template to populate files**

```
{% for host in groups['all'] %}
```

```
    "This is each systems FQDN: {{ hostvars[host]['ansible_fqdn'] }}"
```

```
{% endfor %}
```

Use 'ansible all -m setup > facts' to find names of dictionaries and keys to use in your template files

- **Avoid non-idempotent behavior**
- **Don't use ansible-navigator, use ansible-playbook**
- **Take 3 breaks during exam**
- **Read the "essential information" section very carefully**
- **It doesn't matter how you complete a task, as long as it's completed successfully**

#### **Playbook Commands:**

Description	Command
Check syntax of a Playbook	ansible-playbook --syntax-check
Complete a playbook dry run	ansible-playbook -C
Run a Playbook against a single host (limit)	ansible-playbook <playbook> --limit <host>
Run a Playbook against a list	ansible-playbook -i localhost, server1, server2 <playbook>.yaml
Run a Playbook against a list of hosts	ansible-playbook -i /etc/ansible/hosts <playbook>.yaml
Run a Playbook to update packages on hosts	ansible-playbook all -m command -a "yum update -y"
Which tasks will the playbook run	ansible-playbook <playbook>.yaml --list-tasks
Run a playbook using encrypted password	ansible-playbook <playbook> --ask-vault-pass
Start running a playbook at a specific task	ansible-playbook --step <playbook>.yaml
Start playbook execution as a specific task	ansible-playbook --start-at-task="task name" <playbook>.yaml

--	--

### **Ansible Vault:**

Description	Command
Create an encrypted password file	ansible-vault create/edit/view <vault.yml>
Change password on encrypted file	ansible-vault rekey <vault.yml>
Create ansible vault with a vault ID	ansible-vault create <vault_id>@<vault.yml>
Create directory for encrypted vault and unencrypted vault secrets	mkdir secrets vars # vault goes in vars
Prompt ansible-playbook for vault password	--ask-vault-pass
Point ansible-playbook to unencrypted file containing vault secret	--vault-password-file
Include the vault file with vars_files	vars_files
Use ask vault pass or vault pass file	--vault-password-file OR --ask-vault-pass, not both
Running a playbook with multiple different vaults	ansible-playbook --vault-id dev@dev-password --vault-id prod@prompt site.yml

### **Important Information:**

Description	Command
Ansible Verbosity	ansible-playbook -v -i <HOST> <playbook>.yaml    <-displays output data
	ansible-playbook -vv -i <HOST> <playbook>.yaml    <-displays output & input

	<p>data</p> <p>ansible-playbook -vvv -i &lt;HOST&gt; &lt;playbook&gt;.yaml &lt;-information about connections</p> <p>ansible-playbook -vvvv -i &lt;HOST&gt; &lt;playbook&gt;.yaml &lt;-information about plugins</p>
Ansible Priorities	<p>Priority in which the config files are processed:</p> <ol style="list-style-type: none"> <li>1) ANSIBLE_CONFIG (environment variable)</li> <li>2) ./ansible.cfg (in the current directory)</li> <li>3) ~/.ansible.cfg</li> <li>4) /etc/ansible/ansible.cfg</li> </ol>
Commonly Modified Settings in 'ansible.cfg'	<p>inventory= Location of Ansible inventory file</p> <p>remote user= Remote user account used to establish connections to managed hosts</p> <p>become= Enables/disables privilege escalation for operations on managed hosts (NOT BY DEFAULT!)</p> <p>become_method= Defines privilege escalation method</p> <p>become_user= User account for escalating privileges</p> <p>become_ask_pass= Defines whether privilege escalation prompts for password</p>
Ansible Files (Defaults, if you don't create your own project dirs)	<p>/etc/ansible/hosts # Inventory (hosts) file</p> <p>/etc/ansible/ansible.cfg # Configuration file</p> <p>/etc/ansible/&lt;myplaybook.yml&gt; # Playbook file</p>
Ansible Forks  (Used on devices that don't have a python stack, so processing must be completed on the control node, The default value for forks is 5, and should be increased if managed nodes are	<p>You can manage this setting by using the forks parameter in ansible.cfg or</p> <p>with -f option at command line</p> <p>It prevents the control node from being overloaded</p> <p>Because processing in most environments is done on the managed hosts, the maximum setting of five forks just slows down the working of Ansible, and it is a good idea to increase this maximum to something significantly</p>

linux)	higher. If only Linux hosts are managed, there is no reason to keep the maximum number of simultaneous tasks much lower than 100
<p>Serial Task Execution</p> <p>(By default, Ansible runs task by task. It will run Task 1 on all hosts, then Task 2 on all hosts.</p> <p>If you need it to run only a few hosts at a time, set serial: #)</p>	<p>For instance, the serial: 3 keyword is used in the header of a play, all tasks are executed on three hosts, and after completely running all tasks on three hosts, the next group of three hosts is handled</p> <p>Useful for making sure that not all systems are down at the same time</p> <p>Ex. task: Install updates &amp; reboot</p> <p>With serial: 3 -&gt; All hosts won't be down at the same time, it will run on 3 hosts, then move onto the next 3</p>